

Augmenting Learning Components for Safety in Resource Constrained Autonomous Robots

Shreyas Ramakrishna¹, Abhishek Dubey¹, Matthew Buruss¹, Charles Hartsell¹, Nagabhushan Mahadevan¹, Saideep Nannapaneni², Aaron Laszka³, and Gabor karsai¹

¹ Institute for Software Integrated Systems, Vanderbilt University

² Wichita State University

³ University of Houston



Tel (615) 343-7472 Fax (615) 343-7440
1025 16th Avenue South | Nashville, TN 37212
www.isis.vanderbilt.edu



Outline

- Introduction & Research contribution.
- Testbed – DeepNNCar.
- Our Contributions:
 - Blending Control actions for safety.
 - Resource Management.
 - Confidence Estimator.
 - System Integration.
- Experiments & Results.
- Summary & Ongoing work.

Autonomous Robots

- Learning Enabled Components (LEC) driven autonomous robots have been used in self driving, UAV's, Navigation, and off road obstacle avoidance, etc.
- Recently end-to-end learning approach has gained popularity. (Alpha Go & NVIDIA DAVE-II)
- E2E learning uses data-driven models (Neural Networks) to map sensor inputs to control actions.



Middleware Framework

- 1 Augment safety controllers & blend control actions to improve safety
- 2 Resource management & Task Offloading
- 3 System Level Confidence Estimator

Mediated Perception control approach



End-to-End Learning approach



End-to-end reinforcement learning trained AlphaGo



NVIDIA's DAVE-II performed steering using e2e learning.

Problem with data driven LEC models

- Black box with no safety guarantees.
- Testing and verification is challenging.
- Large amounts of data requirements.

Simplex Architectures (SA)

- SA have been used in Safety critical CPS to improve its assurance.
- Integrates high performance controller with safety controllers and arbitrates between them.
- Arbitration techniques: (1) Linear Matrix Reachability, (2) Hybrid System Reachability.

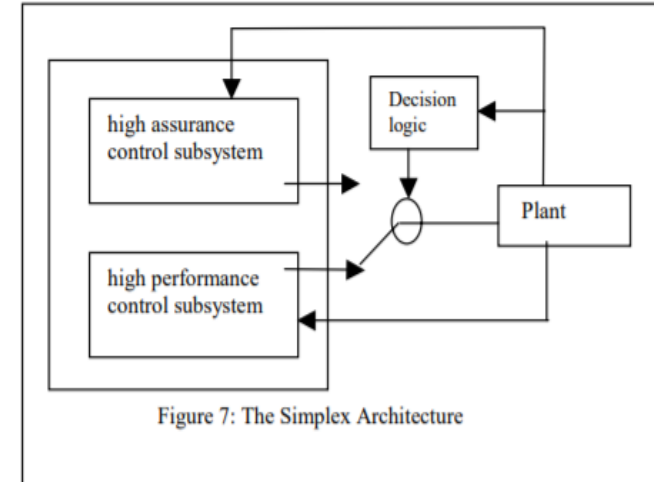
How do we combine the controllers, if we do not have a high assurance controller?

Weighted Simplex Strategy


Idea: Use Weighted Ensemble approach to blend the control actions of the controllers.

$$S_C = w_1 * S_{C1} + w_2 * S_{C2} + \dots + w_N * S_{CN}$$

where, w_1, w_2 are ensemble weights, and S_C = control actions




Simplex Architecture¹



Our controllers do not have high safety assurance

[1] Sha, Lui. "Using simplicity to control complexity." *IEEE Software* 4 (2001): 20-28.

DeepNNCar: Low cost Autonomous Testbed



Sensors:

- (1) Camera capturing images @ 30 FPS with a resolution (320 x 240 x 3).
- (2) Slot type IR Opto-coupler for RPM measurement.

Computing Unit:

Raspberry Pi 3B

Operating modes:


- (1) Data collection, (2) Livestream, and (3) Autonomous.

Acceleration Protocols:

- Default-controlled using Xbox controller.
- Constant throttle.
- Constant throttle adjusted for steering.
- PID controlled cruise control.

Safety Algorithms:

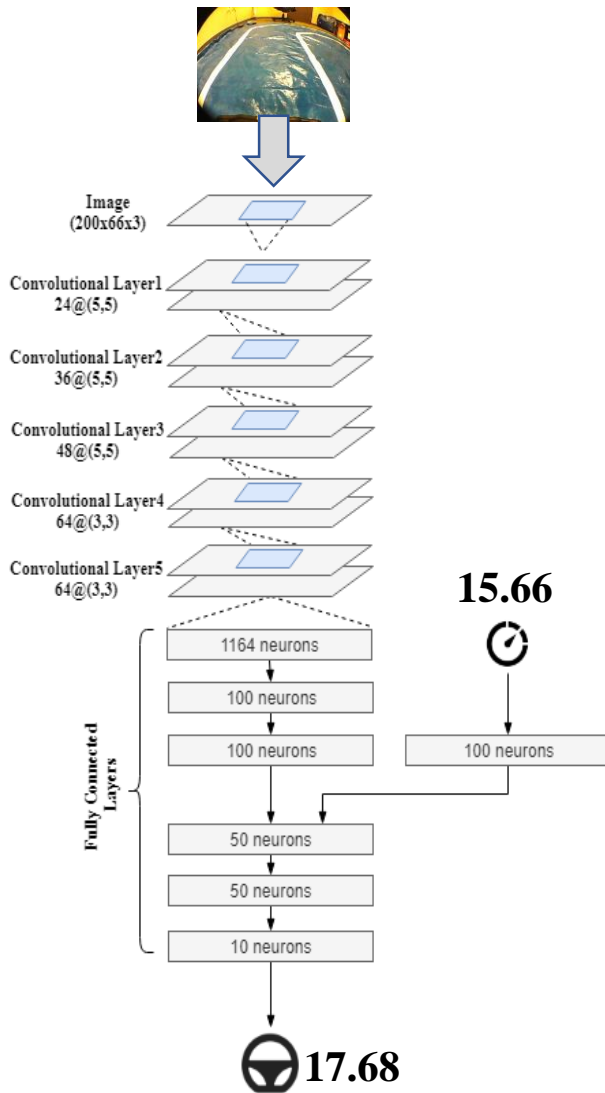
- Lane Detection
- Blur detection



DeepNNCar wirelessly connected to Fog node

github: <https://github.com/scope-lab-vu/deep-nn-car>

CNN for Steering DeepNNCar



Model:

- Modified NVIDIA's DAVE-II¹ Convolutional Neural Network (CNN) to predict steering.
- **Inputs:** Images (200 x 66 x3) and speed (PWM)
- **Output:** Continuous Steering (PWM)

Training:

- 6000 images collected from different tracks.
- Batch size = 128, Epochs = 200, Learning rate = 0.0001

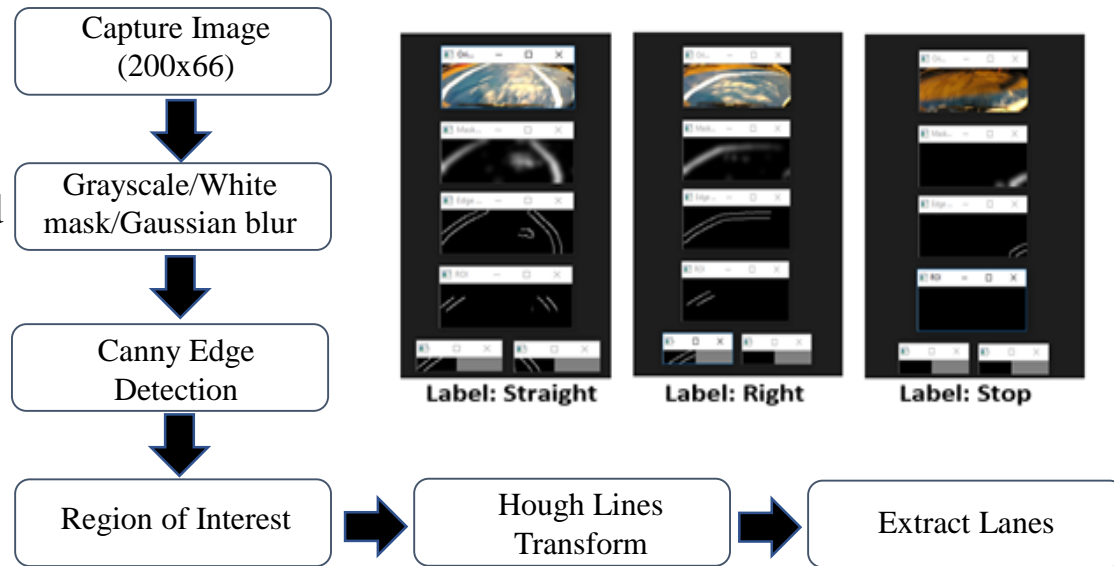
[1] Bojarski, Mariusz, et al. "End to end learning for self-driving cars." *arXiv preprint arXiv:1604.07316* (2016).

Safety Supervisor (SS)

- **Lane Detection (LD) algorithm:** OpenCV based image processing to detect the track segments. (straight, left, right, out)
- Track Segment \longrightarrow Discreet Steering¹
- LD performs with 89.6% accuracy (tested on 3000 images).

Track Segment	Discreet Steering (degrees)
Straight	0
Left	-30
Right	+30
Out of track	NA

Track segment to discreet steering angle conversion

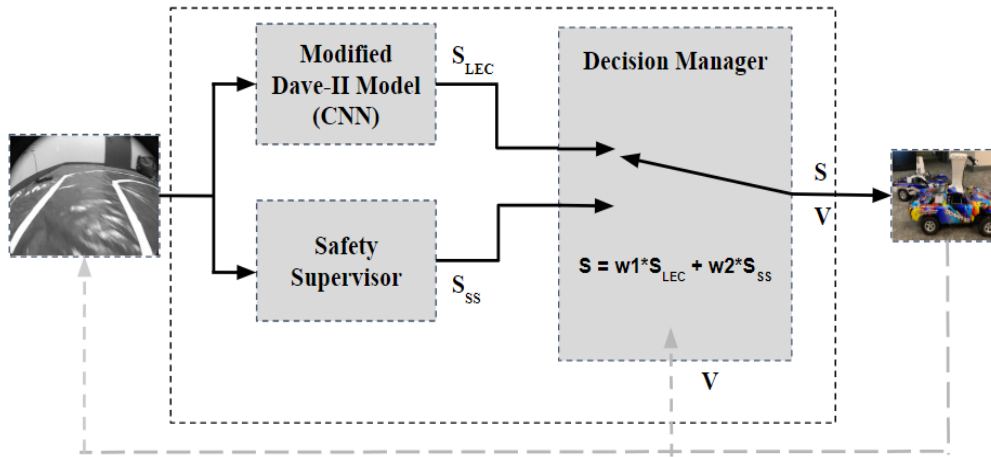


OpenCV image processing steps for Lane Detection

[1] McFall, Kevin. "Using Visual Lane Detection to control steering in a self-driving vehicle." *Smart City 360°*. Springer, Cham, 2016. 861-873.

Blending Control Actions for DeepNNCar

Weighted Simplex Strategy for DeepNNCar



- The steering computation:

$$\text{Steering} = w_{\text{LEC}} * S_{\text{LEC}} + w_{\text{SS}} * S_{\text{SS}}$$

- Speed (V) computation:

$$V = V \pm \Delta V$$

How do we compute the weights?

Can we include domain knowledge (like track segments) in performing the weights calculation?

Also, Literatures from Biswas, Gautham et al.¹ and Yau, Stephen S., et al.², Nascimento, Nathalia, et al. have addressed the importance of using context-sensitive information in different data-driven applications.

[1] Biswas, Gautam, et al. "An approach to mode and anomaly detection with spacecraft telemetry data." *International Journal of Prognostics and Health Management* (2016).

[2] Yau, Stephen S., et al. "Reconfigurable context-sensitive middleware for pervasive computing." *IEEE Pervasive Computing* 1.3 (2002): 33-40.

Simple Weighted Simplex Strategy (SW-Simplex)

- **Idea:** Arguing Machines introduced by Fridman et al¹.
- The control action is decided on the argument: $|S_{\text{LEC}} - S_{\text{SS}}|$.

- If $|S_{\text{LEC}} - S_{\text{SS}}| > T_{\text{ARG}}$ (preset threshold):

$$\text{Steer} = w_1 * S_{\text{LEC}} + w_2 * S_{\text{SS}}.$$

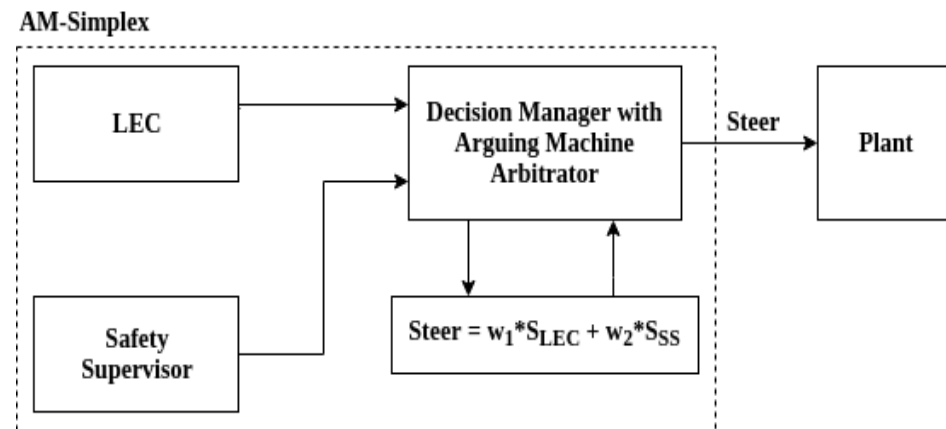
$$\text{Speed} = V_{\text{AM}}$$

- If $|S_{\text{LEC}} - S_{\text{SS}}| \leq T_{\text{ARG}}$

$$\text{Steer} = S_{\text{LEC}}$$

$$\text{Speed} += V_{\text{AM}}$$

- Weights: $W_1 = 0.8$, $W_2 = 0.2$ (after several test runs)



Arguing Machine for DeepNNCar

- **Parameter selection:** trial and error experimental runs.
- **Questions:** Can the weights dynamically change? Can we automate the parameter tuning process?

[1] Fridman, Lex, Benedikt Jenik, and Bryan Reimer. "Arguing machines: Perception control system redundancy and edge case discovery in real-world autonomous driving." *arXiv preprint arXiv:1710.04459* (2017).

Context-Sensitive Weighted Simplex (CSW-Simplex)

- **Reinforcement Learning (RL)**¹ to find context-sensitive dynamically changing weights.
- RL is a type of dynamic programming that trains algorithms using a system of reward and punishment.
- Specifically we use Q-learning² algorithm.

Important: Designing a reward function to: (1) avoid moving off track, and (2) optimize speed.

Q-learning Specifics:

States: $(V, w, 1-w)$ where,

- V is the velocity of the car, $V \in [15.58, 15.62]$
- w is the weight, $w \in [0, 1]$
- The number of states in our case was 102

Actions: $(\Delta V, \Delta w, 1-\Delta w)$ where,

- ΔV is the change in the velocity (possible $\Delta V \in [0, 0.010, -0.010]$)
- Δw is the change in the weights (possible $\Delta w \in [0, 0.005, -0.05]$)

action space	$\uparrow V$ by 0.01	$\downarrow V$ by 0.01	NOP
$\uparrow W_L$ by 0.05	(0.55, 0.45, 15.86)	(0.55, 0.45, 15.84)	(0.55, 0.45, 15.85)
$\downarrow W_L$ by 0.05	(0.45, 0.55, 15.86)	(0.45, 0.55, 15.84)	(0.45, 0.55, 15.85)
NOP	(0.50, 0.50, 15.86)	(0.50, 0.50, 15.84)	(0.50, 0.50, 15.85)

Sample set of actions the car can perform in each state

Q-Learning Setup

Reward: Formulated to account for speed and deviation from track center.

$$R(s_t, a_t) = V - V * \text{trackPos}$$

where, $R(s_t, a_t)$: reward.

V_t : velocity of the car.

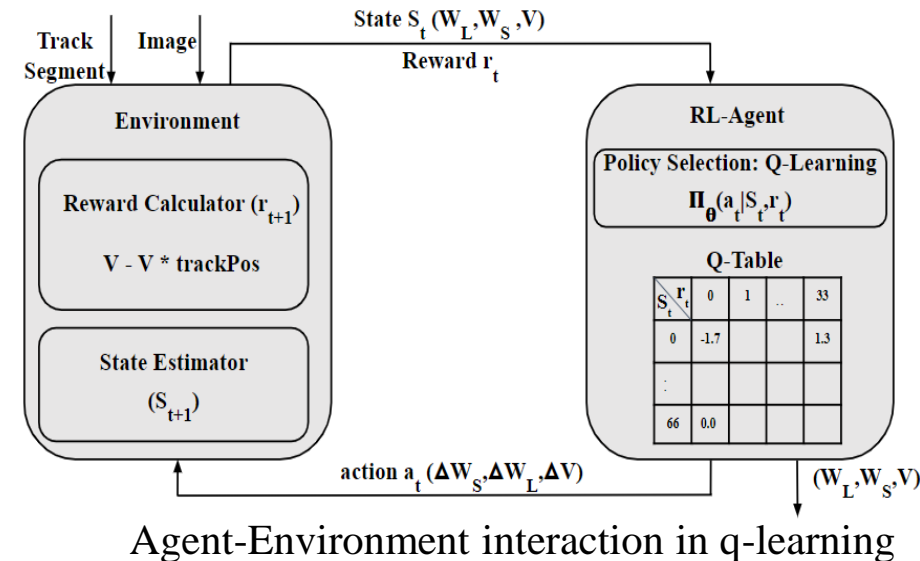
trackPos: scalar value based on position of the car.

- Q-value for (state, action) is computed using Bellman equation.

$$Q_{\text{new}}(s_t, a_t) = Q(s_t, a_t) + \alpha [R(s_t, a_t) + \gamma * \max_{a'} Q'(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

Where, α = learning rate (0.1), γ = discount factor (0.4)



- Designing RL reward function is important to learn driving policy.
- Any RL algorithm can be used to find the weights



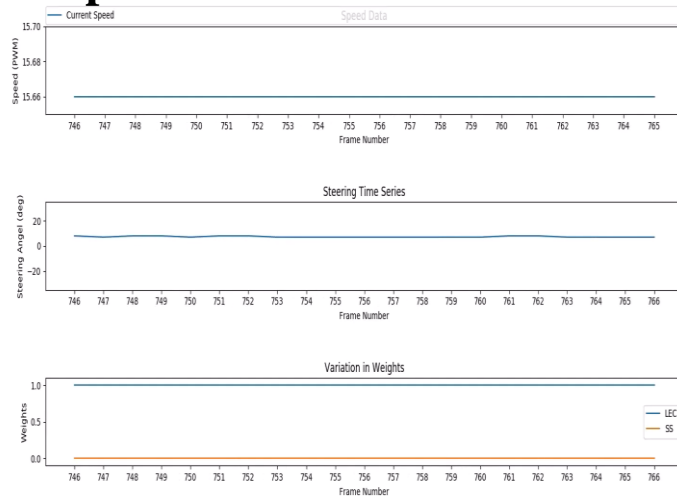
Exploration: Stores state, action and q-value in the Q-table. (1000 steps)

Exploitation: RL-Agent uses the Q-table to choose the actions.


LEC & Safety Supervisor



DeepNNCar Runtime Statistics



DeepNNCar Runtime Statistics






SW-Simplex & CSW-Simplex






Did we create a new problem?

The complex computations of simplex strategies increases the core temperature and CPU utilization of RPi3.

DeepNNCa




Runtime Statistics



Resource Manager and Task Offloader

- Complex computations of simplex strategies increases the temperature and cpu utilization of RPi3.
- Continuously monitor the RPi3 temperature.
- Offload tasks to maintain temperature within limits (70°C).
- Task Offloader: selects the fog with least latency.
- To compensate for increased pipeline times, top speed of the



Top speed calculated as a function of added latency
+ Position of car on track.

Fog Node Selection

FogNode1 (Latency ms)	FogNode2 (Latency ms)	FogNode1 (Average Latency ms)	FogNode2 (Average Latency ms)	Fog Node Selected
13	10.8	11.87	12.17	FogNode1
11.4	11.7			
11.2	14.0			
11.7	11.4	11.23	10.77	FogNode2
11.9	10.6			
10.1	10.3			

Continuous background pings to available fog nodes using iperf every 30 seconds.

Task movement as temperature varied




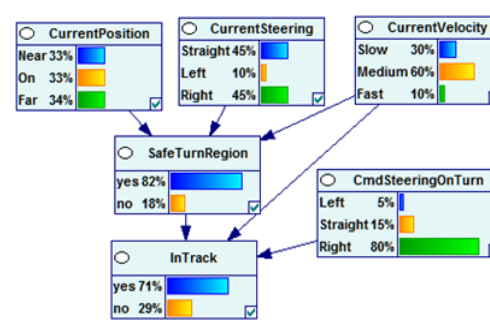
Fig. 12: The effect of offloading the tasks in response to high temperature per iteration of the inference pipeline. The trigger to offload the task is 70°C. The blue line shows the temperature in Celsius. The red line shows when the tasks were offloaded to the fog (on=on fog, off=off fog). The graph shows a subset of iterations (total 10000 iterations).

System Level Confidence Estimation

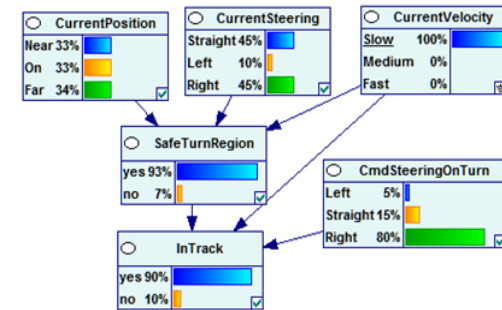
- To monitor the system actions and provide a level of confidence about its safety in different scenarios.
- Bayesian network model to learn a probability distribution about the different scenarios leading to safety of the car.

We use:

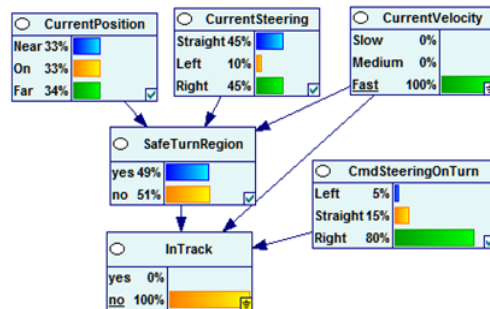
- Car state information: (current-steering, current-velocity, current-position)
- SafeTurnRegion: Calculated from the track, it captures the car being in the non-critical region when it makes the steering decision.



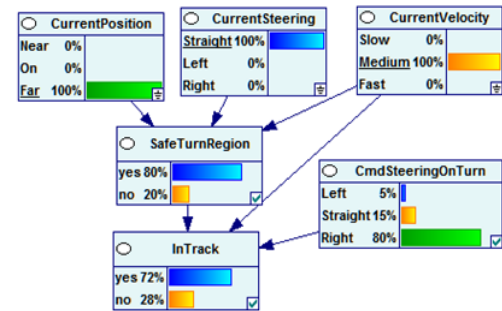
(a) Bayesian Network model for Safety Assurance



(b) Bayesian Inference when current-velocity is set to slow



(c) Bayesian Inference when current-velocity is set to fast

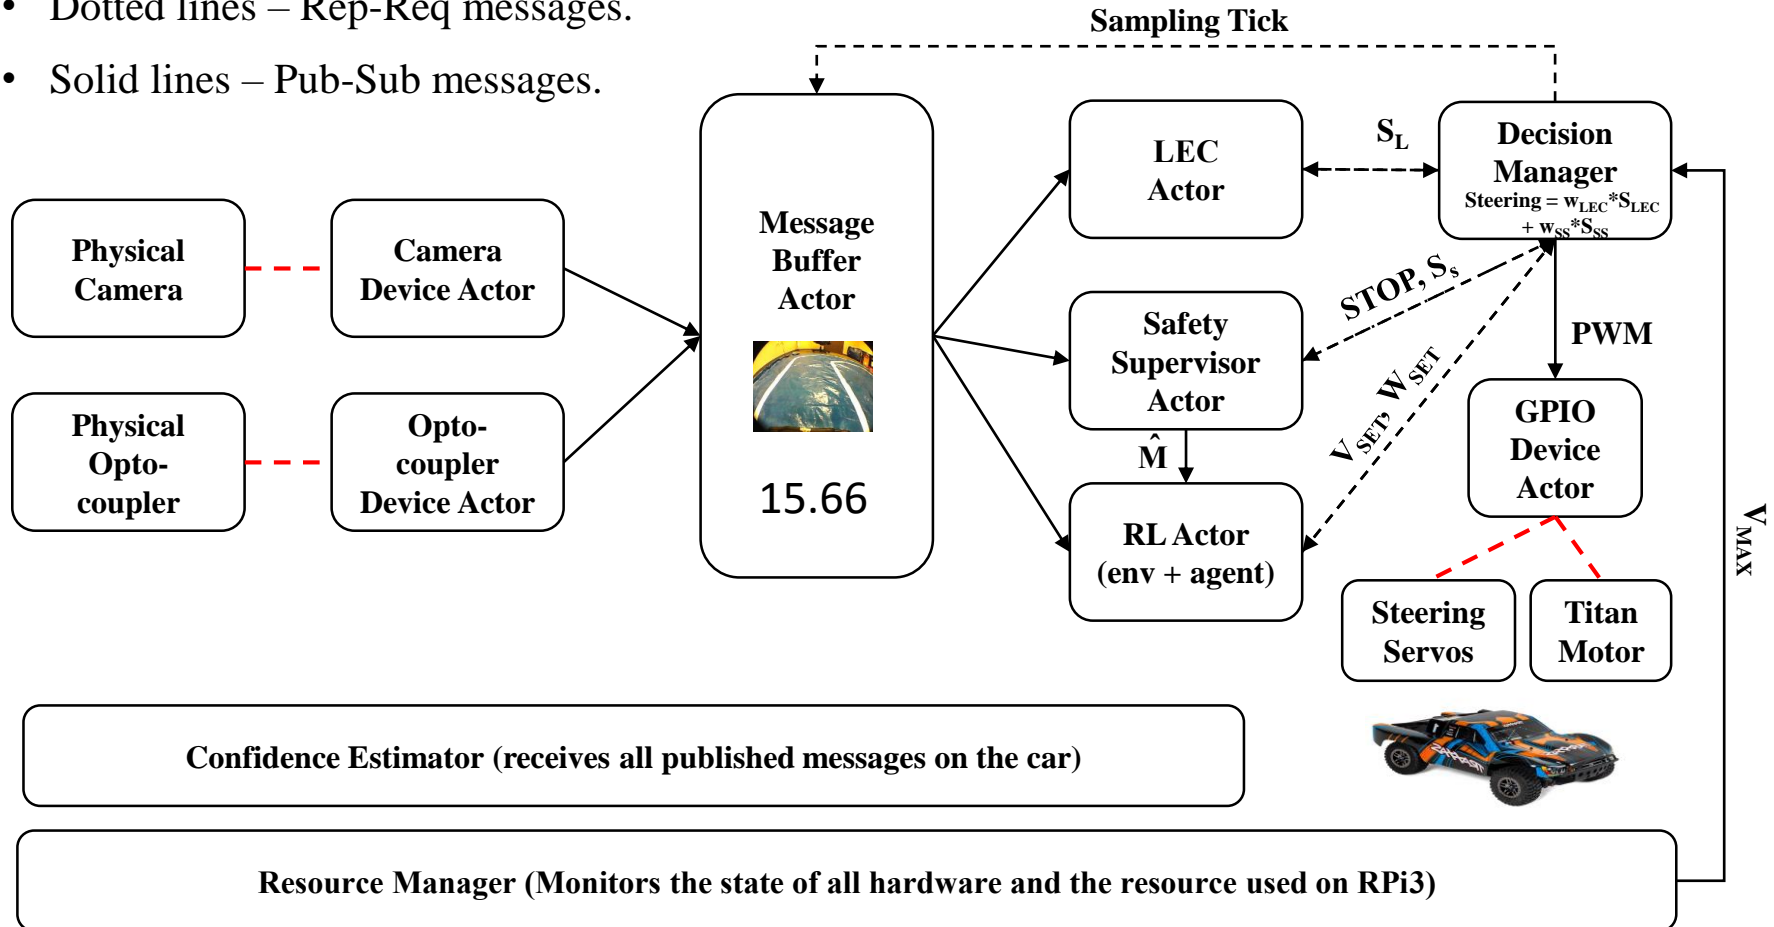


(d) Bayesian Inference when current-velocity is set to medium

The probability of car remaining on track is computed using current position, velocity and steering action.

System Integration

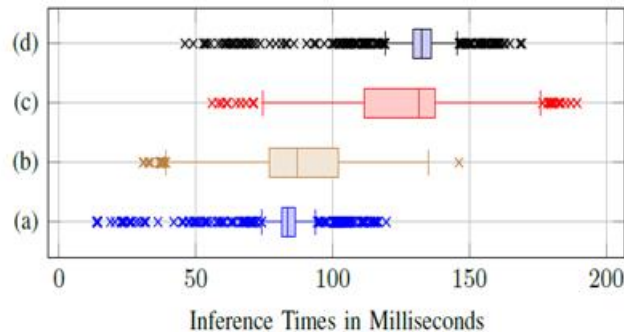
- These components communicate using different messaging patterns of ZeroMQ.
- Red lines – Physical connections.
- Dotted lines – Rep-Req messages.
- Solid lines – Pub-Sub messages.



Performance of weighted simplex strategies

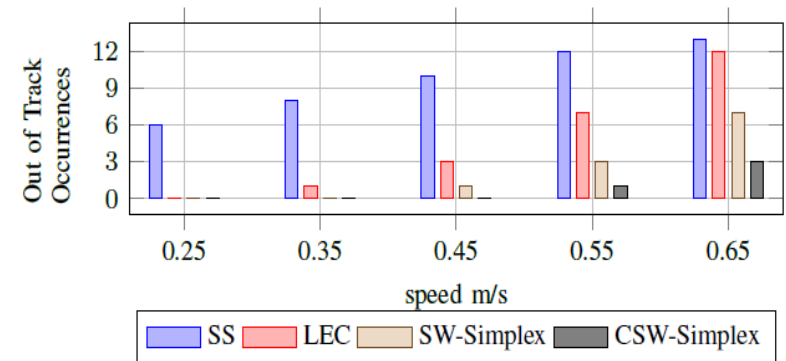


Indoor track designed in our laboratory.




Inference times in milliseconds of (a) Safety Supervisor: driving only with the safety supervisor, (b) LEC: driving only with the modified Dave-II model, (c) SW-Simplex, and (d) CSW-Simplex

SW-Simplex and CSW-Simplex increases the inference pipeline time, due to increased simplex computations.



SW-Simplex and CSW-Simplex reduce safety violations by 40% and 60%.




Speeds in meter per second of (a) Safety Supervisor: Driving only with the safety supervisor, (b) LEC: Driving only with the modified Dave-II model, (c) SW-Simplex, and (d) CSW-Simplex


SW-Simplex and CSW-Simplex optimizes speed (~0.45 m/s) on indoor tracks.

What happens when tasks get offloaded?

Impact of Latency on speed and inference time



Inference times in milliseconds (a) CSW-Simplex with all tasks executed onboard (b) CSW-Simplex with RL-Manager offloaded.




Speed readjustment during offload (m/s) (a) CSW-Simplex with all tasks executed onboard (b) CSW-Simplex with RL-Manager offloaded.

- The impact of added latency reflects as increased pipeline times.

$$T_{\text{Total}} = T_{\text{onboard}} + T_{\text{offload_roundtrip}}$$

T_{onboard} : time for onboard task computation

$T_{\text{offload_roundtrip}}$: round trip time for offloaded tasks



- Top speed is saturated to compensate for increased pipeline time.
- $V_{\text{MAX}} = F(T_{\text{Total}})$, top speed is saturated based on the latency.


- Computations on the fog are ~ 27 times faster.
- However, the varying latency could cause a problem to the safety of the car.

Middleware Framework

- 1 We blend the control actions of controllers using Weighted Simplex Strategies, to improve the systems safety guarantees and optimize its performance.
- 2 We describe a resource management and task offloading scheme to move tasks at runtime to fog nodes.
- 3 We introduce a system level confidence estimator to assess the confidence of the robot systems current action

Ongoing Work

- We are currently extending our work to multi-car platoon setup.
- Ciholas ultra-wideband sensors (CUWB) are used to get the position of the cars at runtime.
- The cars use safety contracts to communicate safe distance information, and the follower cars speed is saturated based on the distance between cars.



Ultra-wideband sensors can be placed around the tracks. Masters, Anchors, and tags communicated to provide Real-time position.



Ultra-wideband sensors attached to a single car.
We get the Real-time position.

<https://cuwb.io/docs/v2.0/overview/>